

Final Project Summary

 **Note:** This summary references historical before→after comparisons (e.g., “8.08 → 3.17 min”) to document the P0 migration improvement. Current production metrics use the spatially-stratified P0 baseline. See [nomenclature_migration.md](#) (nomenclature_migration.md).

EMS Readiness Optimization for Manhattan — v1.3.0

Completion Date: March 15, 2026
Repository: [github.com/cnsp/ems-optimization](#)
Tag: v1.3.0 (Capacity & Baseline Improvements)

Repository Statistics

Metric	Value
Total files in repository	230+
Total Python lines of code	8,500+
Python source files	46
Jupyter notebooks	8
CSV result files	55+
PNG figures	66+
PDF documents	23
Markdown documents	32
LaTeX tables	4
YAML configuration files	5
Unit tests	39 (all passing)

Simulation Statistics

Metric	Value
Total simulation runs	2,700+
Replications per cell	15-30
Simulation duration per run	168 hours (1 week)
Warm-up period per run	24 hours
Total simulated time	~450,000 hours (~51 years)
Total experiments	8 (factorial + CBD + capacity + extended fleet)
Verification tests	4 (all passed)
Validation pilots	3 (all passed)

Key Metrics Achieved (cap=2, K=20)

Metric	P0 (Baseline)	P1 (Demand-Proportional)	P2 (Optimized)
Mean Response Time	3.17 min	2.63 min	2.57 min
P95 Response Time	6.26 min	5.05 min	4.66 min
8-min Coverage	99.6%	99.6%	99.6%
Mean Utilization	7.6%	7.5%	7.5%

Key Observation: P1 (demand-proportional allocation) captures most of the improvement over P0, indicating that matching supply to demand patterns is the primary driver of performance gains. P2's additional optimization yields a further 2.3% reduction in mean response time beyond P1.

Statistical Significance: All results $p < 0.001$. P0 improvement driven by geographic placement.

Phase Completion Status

Phase	Description	Status	Deliverables
Phase 1	Data Processing & EDA	Complete	10 figures, processed data, EDA notebook
Phase 2	Demand & Service Modeling	Complete	NHPP model, lambda tables, 2 spec docs
Phase 3	Optimization Models	Complete	3 MIP models, allocations, comparison
Phase 4	DES Simulation + V&V	Complete	SimPy engine, 39 tests, V&V log
Phase 5	Production Experiments	Complete	1,770 runs, 5 experiments (incl. CBD)
Phase 6	Statistical Analysis	Complete	ANOVA, post-hoc, 5 pub figures, 4 tables
Phase 7	Final Report & Docs	Complete	Technical report, presentation, roadmap
Phase 8	Alternative Analyses	Complete	Distance metric comparison, CBD-focused analysis
Phase 9	Capacity & Baseline Improvements	Complete	Capacity sensitivity (cap 1-5), P0 spatially-stratified, extended fleet analysis

All 9 phases complete. Project delivered.

Files Added in This Session

Part 1: Repository Sync

- Force-added 38 files previously excluded by .gitignore
- Includes all processed data, figures, experiment tables, and maps
- Pushed successfully to origin/main

Part 2: Phase 7 Deliverables

- docs/technical_report.md (.pdf) — 25-page full report
- docs/executive_presentation.md (.pdf) — 10-slide stakeholder deck
- docs/implementation_roadmap.md (.pdf) — 3-phase deployment plan

- docs/project_archive.md (.pdf) — Timeline, decisions, lessons learned
- docs/code_documentation.md (.pdf) — Architecture & API guide
- docs/file_inventory.md (.pdf) — Complete file listing
- README.md — Enhanced with key findings and full documentation
- scripts/generate_summary_dashboard.py — Dashboard generator
- results/figures/project_summary_dashboard.png — Visual summary
- docs/final_summary.md — This file

Part 3: Gap Closure (v1.1.0)

- docs/cbd_definition.md — CBD geographic definition (10 precincts)
- configs/cbd_scenario.yaml — CBD experiment configuration
- scripts/run_cbd_experiment.py — CBD DES experiment runner (330 runs)
- scripts/analyze_queue_metrics.py — Queue metrics analysis script
- scripts/analyze_seasonal_patterns.py — Seasonal pattern analysis script
- notebooks/09_cbd_analysis.ipynb — CBD analysis notebook
- docs/cbd_robustness_analysis.md (.pdf) — CBD findings document
- docs/queue_analysis.md (.pdf) — Queue metrics analysis document
- docs/gap_closure_report.md (.pdf) — Gap closure verification report
- results/simulation/cbd_experiment/ — 330 CBD simulation runs
- results/figures/cbd_*.png — 3 CBD figures
- results/figures/queue_*.png — 4 queue analysis figures
- results/figures/seasonal_*.png — 3 seasonal analysis figures
- results/tables/cbd_*.csv — 2 CBD tables
- results/tables/queue_*.csv — 2 queue tables
- results/tables/seasonal_*.csv — 1 seasonal table

v1.2.0 - Alternative Analyses (March 12, 2026)

New experiments:

- **Distance Metric Comparison:** Haversine vs Manhattan (taxicab) distance — confirmed model robustness; allocations differ at only 2/48 firehouses; simulation performance identical
- **CBD-Focused Optimisation:** Manhattan-wide vs CBD-focused allocation — demonstrated that demand-weighted objective already optimally balances CBD and non-CBD performance; CBD-focused allocation provides no CBD benefit while severely degrading non-CBD service

New deliverables:

- scripts/generate_manhattan_distance_matrix.py — Manhattan distance matrix generation
- scripts/run_distance_comparison_experiment.py — Distance metric comparison experiment
- scripts/run_cbd_focused_optimization.py — CBD-focused optimisation experiment
- data/processed/distance_matrix_firehouse_precinct_manhattan.csv — Manhattan distance matrix
- docs/distance_metric_comparison.md (.pdf) — Distance metric comparison report
- docs/cbd_focused_optimization_analysis.md (.pdf) — CBD-focused analysis report
- docs/alternative_analyses_summary.md (.pdf) — Combined summary of alternative analyses
- results/distance_comparison/ — 4 figures + 2 tables
- results/cbd_focused_comparison/ — 3 figures + 2 tables

Modified files:

- src/ems_readiness/utils/distance.py — Added `manhattan_distance()` function
- src/ems_readiness/optimization/models.py — Added `build_cbd_focused_demand_weighted()` and

`build_cbd_focused_coverage()`

- `docs/technical_report.md` — Added §5.10 and §5.11; updated figures/tables lists

v1.3.0 - Capacity & Baseline Improvements (March 15, 2026)

Phase 9: Major methodological improvements to baseline policy and capacity constraints.

Key changes:

- **Firehouse capacity constraint reduced from 5 to 2** — Full sensitivity analysis (cap 1-5) shows that cap=2 is operationally realistic and optimal for typical fleet sizes
- **Spatially-stratified P0 baseline** — Replaced index-based P0 with latitude-sorted selection ensuring geographic coverage; P0 mean RT improved by 61% (8.08 → 3.17 min at K=20)
- **Extended fleet analysis** — Complete 810-run production experiment with cap=2 and P0 (spatially-stratified) across 9 fleet sizes (K=10-48) × 3 policies × 30 replications
- **38% narrowing of P0-P2 gap** — Spatial stratification shows geographic placement is the dominant factor in EMS response performance

New scripts:

- `scripts/capacity_sensitivity_analysis.py` — Initial cap=2 vs cap=5 comparison
- `scripts/capacity_sensitivity_full_spectrum.py` — Full-spectrum cap 1-5 sensitivity analysis
- `scripts/p0_spatial_analysis.py` — Spatial stratification analysis and visualization
- `scripts/run_production_v2.py` — Extended fleet analysis runner (810 runs)

New results:

- `results/capacity_comparison/` — 60+ files: allocations, figures, tables for cap 1-5
- `results/production_v2/` — Complete extended fleet results: allocations, simulation, figures, tables

New documentation:

- `docs/firehouse_capacity_analysis.md` — Firehouse capacity methodology and initial analysis
- `docs/capacity_sensitivity_analysis.md` — Full-spectrum capacity sensitivity report

Modified files:

- `src/ems_readiness/optimization/policies.py` — Added `spatially_stratified_allocation()` and `spatial_stratification_analysis()`
- `src/ems_readiness/optimization/__init__.py` — Exposed new spatial stratification functions
- `src/ems_readiness/optimization/models.py` — Updated `build_p_median` for capacity-aware allocation; added `solve_model()` convenience function
- `configs/optimization.yaml` — Default capacity changed from 5 to 2
- All documentation files updated to reflect v1.3.0

Project complete. Tagged as v1.3.0 — capacity & baseline improvements with extended fleet analysis results.