

Executive Summary - EMS Readiness Optimization

Bottom Line

Optimized ambulance staging (P2 - Maximal Coverage) reduces average emergency response time by 68% compared to the current uniform deployment, while increasing the proportion of calls answered within 8 minutes from 64% to 95%.

Background

This study used discrete-event simulation to compare three ambulance allocation policies for Manhattan:

Policy	Strategy
P0 - Uniform	Spread ambulances equally across firehouses (current practice proxy)
P1 - Demand-Proportional	Station more units where crashes occur most often
P2 - Maximal Coverage	Optimally position units to maximise coverage (evaluated at 6-min NYC and 8-min NFPA thresholds)

A total of **1,770 simulation replications** were run across five experiment sets (including CBD robustness), using 30 independent replications per scenario with Common Random Numbers (CRN) for fair comparisons.

All analyses use capacity=2 units per firehouse (the operationally optimal default).

Capacity sensitivity analysis (cap 1-5) confirmed that capacity=2 matches or improves upon higher capacity limits at typical fleet sizes; see the technical report §5.12 for details.

Key Findings

1. Response Time Improvement

Metric	P0 (Uniform)	P1 (Demand-Prop.)	P2 (Max Coverage)
Mean Response Time	8.1 min	2.6 min	2.6 min
P90 (90th %ile) RT	19.3 min	5.5 min	5.3 min
6-min Coverage (NYC law)	~50%	~93%	~94%
8-min Coverage (NFPA standard)	64.2%	94.5%	94.8%

All differences are **statistically significant** ($p < 0.001$) with **large effect sizes** (Cohen's $d > 2.0$).

2. Robustness

- **Fleet size:** P2 outperforms P0 at every fleet size tested ($K = 15$ to 40). Even with 25% fewer units ($K = 15$), P2 outperforms P0 at $K = 40$.
- **Demand surges:** Under $2\times$ demand, P2 maintains 85%+ coverage while P0 drops below 50%.
- **Service time variation:** $\pm 20\%$ changes in on-scene service time have minimal impact on the P2 advantage.

3. CBD Robustness

- P2 maintains near-complete coverage (99.3%) even under $2\times$ CBD demand surge
- CBD response times are consistently below 3 minutes for P1 and P2
- Zero queueing across all experiments confirms sufficient fleet capacity

4. Seasonal Stability

- Monthly demand variation is moderate ($CV = 9\%$, amplitude 28%)
- Peak month (October) has only 10.3% more demand than average
- Policy rankings are stable across seasonal variations

5. Statistical Confidence

- All comparisons confirmed by one-way and two-way ANOVA ($\eta^2 > 0.14$ – large effects).
- 95% confidence intervals for response-time differences exclude zero.
- Tukey HSD post-hoc tests confirm all pairwise policy differences.

Recommendation

Implement P2 (Maximal Coverage) allocation as the primary deployment strategy.

Expected Benefits

1. **19% reduction** in average response time (3.17 → 2.57 min at K=20)
2. **Near-perfect 8-minute coverage (NFPA standard)** at 99.6%, with improved **6-minute coverage (NYC law)** reaching ~98%
3. **Robust performance** under demand fluctuations and operational variability
4. **No additional units required** – improvement comes from better positioning

Implementation Considerations

Factor	Detail
Transition	Phase in over 2–4 weeks by shifting units to optimised positions
Technology	Requires a rebalancing algorithm running periodically (e.g., shift changes)
Training	Crews need orientation on new staging locations
Monitoring	Track mean RT, 6-min coverage (NYC law), and 8-min coverage (NFPA standard) weekly to verify gains
Fallback	Maintain ability to revert to demand-proportional (P1) if needed

Cost-Benefit Summary

Item	Estimate
Implementation cost	Low – repositioning existing units, no new hires
Time to deploy	2–4 weeks
Expected lives impacted	Meaningful – faster response is associated with improved survival rates
Risk	Low – P2 holds up across all tested scenarios

Study Limitations

- Travel times use Haversine distance proxy (not real road networks).
- Demand model based on motor-vehicle crashes only; other EMS call types excluded.
- Simulation assumes single-tier response (no differentiated BLS/ALS dispatch).
- Results specific to Manhattan geography and demand patterns.

Next Steps

1. Validate with real FDNY response time data (if obtainable).
 2. Extend to multi-borough deployment.
 3. Incorporate real-time demand forecasting for dynamic rebalancing.
 4. Pilot test with a subset of units in a controlled trial.
-

Analysis performed with 1,770 simulation replications (including CBD robustness), statistical testing (ANOVA, Tukey HSD, Bonferroni corrections), full queueing analysis, seasonal variation assessment, and publication-quality reporting. Full methodology documented in `docs/output_analysis.md` and `docs/technical_report.md`.