

Notebook Guide

Purpose, scope, and recommended usage for the Jupyter notebooks in `notebooks/`.

Overview

The eight notebooks in this project are **exploratory and analytical companions** to the production pipeline, not standalone end-to-end runners. They were created during development to prototype analysis, visualize intermediate results, and document findings for each project phase. The production workflow is driven by the scripts in `scripts/` and the `src/ems_readiness` package.

Bottom line: Use `scripts/` to reproduce results. Use `notebooks/` to understand, explore, and present them.

Notebook Inventory

Notebook	Phase	Purpose	Standalone?
02_eda_spatiotemporal	1 — EDA	Spatiotemporal analysis of crash demand patterns; generates hourly/daily/seasonal figures	Yes (reads raw data directly)
03_input_modeling	2 — Demand	Documents NHPP demand model fitting, lambda tables, goodness-of-fit	Partially (reads processed data)
04_service_travel_proxy	2 — Service	Demonstrates distance matrix, travel-time proxy, service-time distribution, arrival generator	No (imports <code>ems_readiness</code> package)
05_optimization	3 — Optimization	Documents optimization formulations and policy comparison results	Partially (reads results CSVs)
06_simulation_debug	4 — Simulation	Verification and validation of the DES engine (toy examples, pilots)	No (imports <code>ems_readiness</code> package)
07_production_results	5 — Results	Analyzes 1,440 production simulation runs across four experiment sets	Partially (reads results CSVs)
08_statistical_analysis	6 — Statistics	Full statistical analysis: ANOVA, Tukey HSD, effect sizes, confidence intervals	Partially (reads results CSVs)
09_cbd_analysis	Robustness	CBD-focused optimization and distance-metric comparison	Partially (reads results CSVs)

“Standalone?” column explained

- **Yes:** The notebook reads raw data and can run independently given the raw data files.

- **Partially:** The notebook reads pre-computed results (CSVs from `data/processed/` or `results/`). It can run if those files exist, but does not regenerate them — that is the job of `scripts/`.
- **No:** The notebook imports `ems_readiness` as a library (via `sys.path` injection) and depends on the installed or path-accessible package.

Relationship to `scripts/`

<code>scripts/</code>	<code>notebooks/</code>
<code>data_audit.py</code>	-> <code>02_eda_spatiotemporal.ipynb</code>
<code>demand_modeling.py</code>	-> <code>03_input_modeling.ipynb</code>
<code>(package modules)</code>	-> <code>04_service_travel_proxy.ipynb</code>
<code>run_optimization_comparison.py</code>	-> <code>05_optimization.ipynb</code>
<code>run_verification.py</code>	-> <code>06_simulation_debug.ipynb</code>
<code>run_validation_pilots.py</code>	
<code>run_production_experiments.py</code>	-> <code>07_production_results.ipynb</code>
<code>run_production_v2.py</code>	
<code>analyze_production_results.py</code>	-> <code>08_statistical_analysis.ipynb</code>
<code>run_cbd_experiment.py</code>	-> <code>09_cbd_analysis.ipynb</code>
<code>run_cbd_focused_optimization.py</code>	

Key points:

- Scripts generate data and results; notebooks visualize and analyze them.
- Scripts are the authoritative, reproducible pipeline. Notebooks may contain exploratory code that was later formalized into scripts.
- Four notebooks (`05` , `06` , `07` , `08`) have paired `.py` files in `notebooks/` — these are Jupyter text percent-format scripts used for version control of the notebook content.
- No notebook calls a script via `subprocess`. The `08_statistical_analysis` notebook references script paths in documentation but loads results files directly.

Recommended Workflow

To reproduce all results from scratch

```
# 1. Process raw data
make data

# 2. Run the full analysis pipeline
make analysis

# 3. Or run individual scripts
python scripts/run_optimization_comparison.py
python scripts/run_production_v2.py
python scripts/analyze_production_results.py
```

To explore or present results

Open notebooks in JupyterLab after running the production scripts:

```
jupyter lab notebooks/
```

Notebooks assume results files already exist in `data/processed/` and `results/`.

To modify the simulation or optimization

1. Edit source code in `src/ems_readiness/`
2. Re-run the appropriate script in `scripts/`
3. Open the corresponding notebook to visualize updated results

Jupytertext `.py` Files

The `.py` files in `notebooks/` (e.g., `05_optimization.py`) are **Jupytertext percent-format** mirrors of the corresponding `.ipynb` files. They exist for cleaner version control (diffs on `.py` are more readable than on `.ipynb` JSON). These are not independent scripts — they are an alternative representation of the same notebook content.

Maintenance Status

All notebooks reflect the final project state (V2 with spatially-stratified P0, capacity=2 default, CBD robustness analysis). They are maintained as documentation and presentation artifacts. For any future model changes, update the scripts first, re-run experiments, then refresh notebook outputs.