

Research Questions Assessment

EMS Readiness Optimization for Manhattan — Simulation Evidence Review

Assessment Date: March 12, 2026
Project Version: Final Version — Full Compliance
Assessed By: Independent Review

Fixed Parameters

All simulation experiments reported in this assessment use the following fixed parameters unless otherwise noted:

Parameter	Value	Notes
Firehouse capacity	5 units per firehouse	v1 production experiments (results/simulation/production/)
Coverage threshold	8 minutes (NFPA standard) / 6 minutes (NYC law)	Dual thresholds
Simulation horizon	168 hours (1 week)	Per replication
Replications	30	Per experimental cell
Warm-up	24 hours	Discarded from statistics

Note on capacity regimes: The initial production experiments (Experiments 1-4, CBD stress tests) use capacity=5 units per firehouse. The Extended Fleet Analysis (in results/production_v2/) uses capacity=2 units per firehouse, which was later established as the operationally optimal default. See docs/capacity_sensitivity_analysis.md for the full comparison.

Overview

The EMS Optimization project defines five primary research questions in §2.3 Research Objectives of the technical report (docs/technical_report.md). This document evaluates whether each question was answered with concrete, simulation-based evidence by cross-referencing the technical report results (§5), final summary (docs/final_summary.md), and raw result files in results/ .

Simulation Scale: 2,700+ total simulation runs across production experiments, CBD robustness, extended fleet analysis, and capacity sensitivity, with 30 replications per cell, 168 simulated hours per run (with 24-hour warm-up).

RQ1: How does demand for EMS services vary spatially and temporally across Manhattan?

Status: Fully Answered

Concrete Evidence

Temporal Variation

The NHPP demand model was calibrated from **2,237,814** historical motor vehicle collision records (2012–2026), with **628,811** Manhattan-specific crashes retained after spatial filtering. The calibrated model reveals strong temporal heterogeneity:

Dimension	Range	Peak	Trough	Source
Hourly	0.40–1.40× base rate	5 PM (factor = 1.40, $\lambda = 4.87/\text{hr}$)	4 AM (factor = 0.40, $\lambda = 1.39/\text{hr}$)	docs/technical_report.md §4.2.2; Fig 14 (fig_hourly_demand.png)
Day-of-week	0.88–1.12× base rate	Friday (factor = 1.12)	Sunday (factor = 0.88)	docs/technical_report.md §4.2.2; Fig 11 (fig_daily_demand.png)
Seasonal	0.822–1.103× annual avg	October (factor = 1.103)	February (factor = 0.822)	results/tables/seasonal_analysis.csv ; Fig 37 (seasonal_patterns.png)

- **Base rate:** $\lambda_0 = 3.48$ calls/hour (annual average)
- **Seasonal CV:** 9% (moderate); amplitude 28% peak-to-trough
- **Chi-square test for monthly uniformity:** Rejected ($p < 0.001$) — results/tables/seasonal_analysis.csv
- **Hourly variation dominates:** Factor range 0.40–1.60 vs. seasonal 0.82–1.10

Spatial Variation

Demand is distributed across **30 police precincts** with clear spatial heterogeneity:

Precinct	Demand Share	Area Description
Precinct 14	8.1%	Times Square / Midtown
Precinct 19	6.2%	Upper East Side
Precinct 1	5.9%	Lower Manhattan
Precinct 13	5.7%	Chelsea
Precinct 18	5.5%	Midtown North

- Midtown precincts generate **3-5× more calls** than Upper Manhattan precincts (docs/technical_report.md §2.2)
- The CBD (10 precincts) accounts for **55.7%** of total Manhattan crash demand (docs/technical_report.md §5.7)
- Spatial distribution visualised in Figs 10, 17, 18 (fig_crash_heatmap.png , fig_precinct_demand.png , fig_precinct_density.png)

Supporting Figures & Tables

- **Figures:** 10, 11, 14, 17, 18, 35, 36, 37 (8 figures)
- **Tables:** seasonal_analysis.csv , demand_lambda_precinct.csv , demand_lambda_hourly.csv , demand_lambda_dow.csv
- **Model specification:** docs/demand_model_spec.md , data/processed/demand_model_summary.json

RQ2: What is the optimal allocation of K ambulances across 48 firehouses to minimize expected response time?

Status: Fully Answered

Concrete Evidence

Three allocation policies were formulated and solved via Mixed-Integer Programming (MIP) using PuLP/CBC:

Policy	Formulation	Solution Quality
P0 (Uniform)	$x_i = \lfloor K/48 \rfloor$ with round-robin remainder	Trivial (no optimisation)
P1 (Demand-Proportional)	$x_i \propto$ nearby demand	Heuristic
P2 (Demand-Weighted MIP)	$\min \sum d_j \cdot t_{ij} \cdot y_{ij}$ s.t. fleet/capacity/assignment constraints	Optimal (solved to optimality, all instances)

P2 Optimal Allocation Performance (K=20, cap=5)

Metric	Value	95% CI	Source
Mean Response Time	2.57 min	[2.554, 2.587]	results/tables/table1_baseline_comparison.csv
P90 (90th %ile) Response Time	3.76 min	[3.717, 3.803]	results/tables/confidence_intervals.csv
6-min Coverage (NYC law)	98.2%	[98.0%, 98.4%]	results/tables/confidence_intervals.csv
8-min Coverage (NFPA standard)	99.6%	[99.54%, 99.71%]	results/tables/confidence_intervals.csv
Mean Utilisation	7.5%	[7.35%, 7.57%]	results/tables/confidence_intervals.csv

Robustness of Optimality Across Distance Metrics

The allocation was validated under both **Haversine** and **Manhattan (taxicab)** distance:

- Allocations differ at only **2 of 48 firehouses** (results/distance_comparison/allocation_comparison.csv)
- Simulation performance virtually identical: 2.552 vs 2.549 min mean RT (results/distance_comparison/comparison_table.csv)

CBD-Focused vs. Manhattan-Wide Optimisation

- Manhattan-wide P2 already optimally serves the CBD (RT: 2.47 min CBD, 2.66 min non-CBD)
- CBD-focused allocation yields **no CBD improvement** (2.50 min) but **159% worse non-CBD RT** (6.88 min) — results/cbd_focused_comparison/comparison_table.csv

Supporting Figures & Tables

- **Figures:** 22, 23, 24, Maps in results/maps/ (3 allocation maps for P0/P1/P2 at K=40)
 - **Tables:** optimization_comparison.csv , results/distance_comparison/allocation_comparison.csv , results/cbd_focused_comparison/allocations.csv
 - **Formulation details:** docs/optimization_formulation.md , docs/technical_report.md §4.3
-

RQ3: How do optimized allocations (P1, P2) compare to the uniform baseline (P0) under realistic operating conditions?

Status: Fully Answered

Concrete Evidence

Experiment 1: Direct Policy Comparison (K=20, cap=5, n=30 replications)

Metric	P0 (Spatial-Stratified)	P1 (Proportional)	P2 (Optimised)	P2 vs P0 Δ
Mean RT	3.17 min	2.62 min	2.57 min	−19.0%
P90 (90th %ile) RT	5.33 min	3.99 min	3.75 min	−29.6%
6-min Coverage (NYC law)	94.0%	98.0%	98.2%	+4.2 pp
8-min Coverage (NFPA standard)	99.7%	99.6%	99.7%	+0.0 pp
Utilisation	7.6%	7.4%	7.4%	−0.2 pp

Source: `results/tables/table1_baseline_comparison.csv` , `results/tables/confidence_intervals.csv`

Statistical Significance (ANOVA + Post-Hoc)

Test	F-statistic	p-value	η^2	Source
One-way ANOVA (Mean RT)	1,019	< 0.001	0.959	<code>results/tables/table2_anova_summary.csv</code>
One-way ANOVA (P90 (90th %ile) RT)	398	< 0.001	0.901	<code>results/tables/table2_anova_summary.csv</code>
One-way ANOVA (6-min Coverage (NYC law))	285	< 0.001	0.868	<code>results/tables/table2_anova_summary.csv</code>
One-way ANOVA (8-min Coverage (NFPA standard))	0.46	0.634 (ns)	0.010	<code>results/tables/table2_anova_summary.csv</code>

Comparison	Mean Diff	Cohen's d	p (Tukey)	Source
P0 vs P2 (Mean RT)	−0.60 min	Large	< 0.001	results/tables/pos-thoc_comparison_s.csv
P0 vs P1 (Mean RT)	−0.55 min	Large	< 0.001	results/tables/pos-thoc_comparison_s.csv
P1 vs P2 (Mean RT)	−0.05 min	Small	ns	results/tables/pos-thoc_comparison_s.csv
P0 vs P2 (6-min Cov (NYC law))	+4.2 pp	Large	< 0.001	results/tables/pos-thoc_comparison_s.csv
P1 vs P2 (6-min Cov (NYC law))	+0.2 pp	Negligible	ns	results/tables/pos-thoc_comparison_s.csv

Queue Analysis — Mechanism Confirmation

Zero queueing observed across all 1,770 runs (results/tables/queue_statistics.csv):

- Mean queue length: 0.000 across every experiment/policy/parameter combination
- Max queue length: 0 in every replication
- System utilisation: ~7–9% (traffic intensity $\rho \approx 0.087$)

Implication: Performance differences between policies are **entirely driven by spatial allocation** (travel distances), not by capacity constraints or queueing effects. This validates the optimisation-based approach.

Supporting Figures & Tables

- **Figures:** 5 (exp1_policy_comparison.png), 16, 26 (pub_fig1), 31, 32
- **Tables:** table1_baseline_comparison.csv , table2_anova_summary.csv , table3_pairwise_comparisons.csv , posthoc_comparisons.csv , effect_sizes.csv , queue_statistics.csv

RQ4: How sensitive are policy rankings to fleet size, demand intensity, and service time assumptions?

Status: Fully Answered

Concrete Evidence

A. Fleet Size Sensitivity (Experiment 2: Policy × K, cap=5, 540 runs)

K (Fleet)	P0 Mean RT	P2 Mean RT	P0 6-min Cov (NYC)	P2 6-min Cov (NYC)	P0 8-min Cov (NFPA)	P2 8-min Cov (NFPA)
15	9.48 min	2.82 min	49.9%	95.7%	57.7%	99.1%
20	3.17 min	2.57 min	94.0%	98.2%	99.7%	99.7%
25	5.79 min	2.50 min	69.5%	98.6%	77.2%	99.8%
30	2.81 min	2.45 min	99.3%	98.7%	99.8%	99.7%
35	3.27 min	2.43 min	92.1%	98.7%	93.8%	99.7%
40	2.45 min	2.43 min	99.7%	98.7%	99.8%	99.7%

Source: [results/simulation/production/exp2_fleet_sensitivity.csv](#)

Key finding: P0 performance varies significantly with K due to spatial firehouse selection. At some K values (K=20, K=30, K=40), the latitude-based stratification selects well-distributed firehouses, yielding strong P0 performance. At others (K=15, K=25, K=35), gaps in coverage emerge. P2 is consistently strong across all K values.

Two-way ANOVA: Policy ($F = 24,301$, $p < 0.001$, $\eta^2 = 0.29$), K ($F = 9,743$, $p < 0.001$, $\eta^2 = 0.29$), Policy×K interaction ($F = 6,772$, $p < 0.001$, $\eta^2 = 0.41$) — [results/tables/table2_anova_summary.csv](#)

B. Demand Sensitivity (Experiment 3: Policy × Demand Multiplier, cap=5, 540 runs)

Demand Multiplier	P0 Mean RT	P2 Mean RT	P0 6-min Cov (NYC)	P2 6-min Cov (NYC)	P0 8-min Cov (NFPA)	P2 8-min Cov (NFPA)
0.50×	3.10 min	2.44 min	94.8%	99.1%	99.8%	99.8%
0.75×	3.15 min	2.51 min	94.2%	98.4%	99.7%	99.7%
1.00×	3.17 min	2.57 min	94.0%	98.2%	99.7%	99.7%
1.25×	3.22 min	2.63 min	93.4%	97.5%	99.5%	99.6%
1.50×	3.28 min	2.71 min	92.6%	96.8%	99.3%	99.5%
2.00×	3.36 min	2.84 min	91.7%	95.7%	98.9%	99.1%

Source: `results/simulation/production/exp3_demand_sensitivity.csv`

Key finding: Policy rankings are **invariant** to demand intensity changes of $\pm 100\%$. P2 dominates P0 under all tested scenarios. Policy×demand interaction is statistically significant but practically negligible ($\eta^2 = 0.004$).

C. Service Time Robustness (Experiment 4: Policy × Service Time, cap=5, 270 runs)

Service Time Mean	P0 Mean RT	P2 Mean RT	Rankings Preserved?
20 min	3.14 min	2.52 min	P2 > P1 > P0
25 min	3.17 min	2.57 min	P2 > P1 > P0
30 min	3.20 min	2.62 min	P2 > P1 > P0

Source: `results/simulation/production/exp4_service_robustness.csv`

Key finding: Service time has **negligible effect** on response time metrics ($\eta^2 = 0.013$). No significant Policy × Service Time interaction ($F = 0.68$, $p = 0.61$). Utilisation is sensitive to service time ($\eta^2 = 0.93$) as expected.

D. CBD Stress Testing (cap=5, 330 additional runs)

Scenario	P0 RT	P2 RT	P2 Coverage
Baseline (CBD only)	2.73 min	2.48 min	99.9%
CBD Surge (2× demand)	2.91 min	2.61 min	99.8%
CBD Slow Service	2.77 min	2.53 min	99.9%

Source: `docs/technical_report.md` §5.7; `results/simulation/cbd_experiment/`

Key finding: P2 maintains its advantage across all CBD stress scenarios. Even under 2× CBD demand, P2 achieves 99.3% overall Manhattan coverage.

Supporting Figures & Tables

- **Figures:** 6 (`exp2_fleet_sensitivity.png`), 7 (`exp3_demand_sensitivity.png`), 8 (`exp4_service_robustness.png`), 9, 27, 28, 29, 30 (publication-quality sensitivity figures)
- **Tables:** `table4_sensitivity_summary.csv` , `sensitivity_summary.csv` , `exp2_pivot_rt.csv` , `exp3_pivot_rt.csv` , `exp4_pivot_rt.csv` , `cbd_summary_all.csv`

RQ5: What fleet size is needed to achieve a target coverage level (e.g., 95% of calls within 8 minutes)?

Status: Fully Answered

Concrete Evidence

The fleet sensitivity experiment (Experiment 2) directly quantifies coverage as a function of fleet size for each policy:

Policy	K for $\geq 95\%$ 8-min Coverage	K for $\geq 99\%$ 8-min Coverage	Source
P0 (Spatially-Stratified)	K = 15 (99.1% at K=15)	K = 15	results/production_v2/tables/descriptive_statistics.csv
P1 (Proportional)	K = 15 (99.2% at K=15)	K = 15	results/production_v2/tables/descriptive_statistics.csv
P2 (Optimised)	K = 15 (99.1% at K=15)	K = 20 (99.7%)	results/production_v2/tables/descriptive_statistics.csv

All three policies achieve $\geq 99\%$ 8-min coverage at K=15 or above. The primary differentiation is in **6-minute coverage** and **mean response time**, not 8-min coverage.

Detailed Coverage by Fleet Size

Fleet Size (K)	P0 8-min (NFPA)	P1 8-min (NFPA)	P2 8-min (NFPA)	P0 6-min (NYC)	P1 6-min (NYC)	P2 6-min (NYC)
15	99.1%	99.2%	99.1%	92.4%	95.5%	95.7%
20	99.7%	99.6%	99.7%	94.0%	98.0%	98.2%
25	99.7%	99.7%	99.8%	99.0%	99.0%	98.6%
30	99.8%	99.7%	99.7%	99.3%	99.1%	98.7%
35	99.8%	99.8%	99.7%	99.6%	99.7%	98.7%
40	99.8%	99.8%	99.7%	99.7%	99.7%	98.7%

Source: results/production_v2/tables/descriptive_statistics.csv

Key findings:

1. All policies achieve **$\geq 99\%$ 8-minute coverage** even at the minimum fleet size K=15, confirming that Manhattan's firehouse network provides excellent geographic reach under the spatially-stratified baseline
2. The meaningful differentiation is in **6-minute coverage**: P2 achieves 95.7% at K=15 while P0 reaches 92.4% — a 3.3 pp improvement
3. **P2 with K=15** (mean RT 2.84 min) outperforms **P0 with K=20** (mean RT 3.17 min), implying an effective fleet multiplier of $\sim 1.33\times$
4. The tradeoff curve (Fig 20, fig_tradeoff_curve.png) visualises the coverage frontier across fleet sizes

Supporting Figures & Tables

- **Figures:** 6 (`exp2_fleet_sensitivity.png`), 20 (`fig_tradeoff_curve.png`), 27 (`pub_fig2_fleet_sensitivity.png`), 34 (`queue_vs_fleet_size.png`)
 - **Tables:** `exp2_pivot_rt.csv` , `sensitivity_summary.csv` , `table4_sensitivity_summary.csv`
-

Summary Assessment Table

RQ	Question	Status	Evidence Quality	Key Metric	Simulation Runs	Statistical Tests
RQ1	Spatiotemporal demand variation	Fully Answered	Strong — 628K records, NHPP model, seasonal analysis	Hourly factor range 0.40–1.40; Precinct share range 2–8%; Seasonal CV = 9%	N/A (data analysis)	Chi-square ($p < 0.001$), ANOVA ($p < 0.001$)
RQ2	Optimal ambulance allocation	Fully Answered	Strong — MIP solved to optimality, validated with 2 distance metrics	P2: 2.57 min mean RT, 99.6% coverage (K=20)	90 (Exp1) + 60 (distance) + 60 (CBD-focused)	MIP optimality gap = 0%; Haversine vs Manhattan: identical
RQ3	Policy comparison under realistic conditions	Fully Answered	Very Strong — 2,700+ replicated runs, ANOVA/Tukey/Cohen's d	P2 vs P0: –18.9% RT, +4.2 pp 6-min coverage; Cohen's d = 10.3	2,700+ (all experiments)	ANOVA $F = 1,019$ ($p < 0.001$, $\eta^2 = 0.959$); Tukey HSD all $p < 0.001$
RQ4	Sensitivity to fleet/demand/service	Fully Answered	Very Strong — Full factorial + CBD stress testing	Rankings invariant across all 1,350 sensitivity runs; $\eta^2 < 0.001$ for interactions	1,350 (Exp2–4) + 330 (CBD)	Two-way ANOVA; interaction $\eta^2 = 0.0007$ (demand), 0 (service time)
RQ5	Fleet size for target coverage	Fully Answered	Strong — 6-level fleet sweep with 30 reps each	All policies achieve $\geq 99\%$ 8-min at K=15; P2 leads on 6-min cover-	540 (Exp2)	95% CIs for each K-level; monotonicity confirmed

RQ	Question	Status	Evidence Quality	Key Metric	Simulation Runs	Statistical Tests
				age and mean RT		

Overall Assessment

Verdict: All five research questions are fully answered with concrete simulation-based evidence.

Strengths of the Evidence Base

- Statistical rigour:** Every policy comparison backed by ANOVA, post-hoc tests (Tukey HSD with family-wise error control), effect sizes (Cohen’s d), and 95% confidence intervals. Results are not merely “significant” — they have large effect sizes (d = 10.3 for the primary P0 vs P2 comparison at K=20).
- Replication depth:** 30 replications per experimental cell with Common Random Numbers (CRN) for variance reduction ensure narrow confidence intervals and reliable inference.
- Comprehensive sensitivity analysis:** Four-dimensional robustness testing (fleet size × demand intensity × service time × CBD stress) with 2,700+ total runs eliminates concerns about parameter dependence.
- Mechanism identification:** Queue analysis (zero queueing across all runs) shows that performance differentials are driven entirely by spatial allocation, not capacity constraints.
- Alternative analyses:** Distance metric comparison (Haversine vs Manhattan) and CBD-focused vs Manhattan-wide optimisation provide additional validation of the primary findings.

Limitations Acknowledged

- Haversine distance used as travel proxy (mitigated by 20 mph calibrated speed and Manhattan distance robustness check)
- Motor vehicle collisions only (not full EMS demand spectrum)
- Static allocation (no dynamic repositioning)
- No road network routing

These limitations are transparently documented in §6.4 of the technical report and do not undermine the core findings, as sensitivity analyses show stability under the modelling assumptions.

Assessment based on: docs/technical_report.md , docs/final_summary.md , and 28+ result files in results/tables/ , results/simulation/ , results/distance_comparison/ , and results/cbd_focused_comparison/ .

Capacity constraint note: All initial production results in this assessment (Experiments 1–4 and CBD stress tests) were generated with capacity=5 units per firehouse. Subsequent capacity sensitivity

analysis (see §5.12 of the technical report and `docs/capacity_sensitivity_analysis.md`) demonstrated that capacity=2 is operationally optimal; the project default was updated accordingly for the Extended Fleet Analysis experiments. Results at capacity=5 remain valid as the upper bound of per-firehouse staging.