

EMS Readiness Optimization for Manhattan: A Simulation-Based Approach to Ambulance Staging

Comprehensive Technical Report

Authors: EMS Optimization Research Team

Date: March 12, 2026

Version: 2.0.0 (Phase 21 · Full Compliance)

Abstract

This study presents a simulation-based optimization framework for strategic ambulance staging across 48 FDNY firehouses in Manhattan, New York City. Using 2.24 million historical motor vehicle collision (MVC) records from NYC Open Data (2012-2026), we calibrate a Non-Homogeneous Poisson Process (NHPP) demand model with hourly and day-of-week intensity factors (base rate $\lambda = 3.48$ calls/hour). Three allocation policies are evaluated: a uniform baseline (P0), a demand-proportional heuristic (P1), and a demand-weighted Mixed-Integer Programming (MIP) optimized allocation (P2). A discrete-event simulation (DES) engine built with SimPy executes 1,770 production runs across five experiment sets including policy comparison, fleet sensitivity, demand sensitivity, service robustness, and CBD-focused stress tests with 30 replications each using Common Random Numbers for variance reduction. Results demonstrate that the optimized policy P2 reduces mean response time by 68.2% (from 8.08 to 2.57 minutes) and improves 8-minute coverage from 64.4% to 99.6% compared to the uniform baseline, with all differences statistically significant ($p < 0.001$, Cohen's $d > 28$). Performance gains are robust to demand fluctuations ($0.5 \times$ to $2.0 \times$ multiplier), service time variations (20-30 min mean), and CBD-specific surge scenarios. Queue analysis confirms zero waiting across all experiments, establishing that performance differentials arise entirely from spatial allocation efficiency. We recommend adoption of P2 for operational deployment, with a phased 12-month implementation roadmap targeting the highest-impact firehouses first.

Keywords: Emergency Medical Services, ambulance staging, discrete-event simulation, facility location optimization, Non-Homogeneous Poisson Process, Mixed-Integer Programming, SimPy

Table of Contents

1	Executive Summary
2	Introduction
3	Literature Review
4	Methodology
5	Results
6	Discussion
7	Conclusions and Recommendations
8	References
9	Appendices
10	List of Figures
11	List of Tables
12	Reproducibility

1. Executive Summary

Problem Statement

Emergency Medical Services (EMS) in Manhattan face a critical challenge: the current uniform ambulance allocation policy (P0) distributes units equally across 48 FDNY firehouses without regard to spatiotemporal demand patterns, resulting in suboptimal response times and inadequate coverage. With approximately 3.48 motor vehicle collision (MVC) calls per hour and significant variation by time-of-day and precinct, a more intelligent allocation strategy is needed to improve emergency response performance.

Methodology Overview

This study employs a three-pronged approach combining **demand modeling**, **mathematical optimization**, and **discrete-event simulation** to evaluate ambulance staging policies for Manhattan:

- 1 **Non-Homogeneous Poisson Process (NHPP)** demand model calibrated from 2.24 million historical MVC records
- 2 **Mixed-Integer Programming (MIP)** optimization models generating three allocation policies:
 - **P0** (Uniform): Equal distribution across all firehouses
 - **P1** (Demand-Proportional): Units allocated proportional to nearby demand
 - **P2** (Demand-Weighted Optimized): MIP-optimized allocation minimizing expected response time
- 3 **Discrete-Event Simulation (DES)** with 1,770 production runs across 5 experiment sets (including CBD robustness)

Key Findings

Metric	P0 (Current)	P1 (Proportional)	P2 (Optimized)	P2 vs P0 Improvement
Mean Response Time	8.08 min	2.63 min	2.57 min	-68.2%
P90 Response Time	19.47 min	4.03 min	3.76 min	-80.7%
8-min Coverage	64.4%	99.6%	99.6%	+35.2 pp
Mean Utilization	9.1%	7.5%	7.5%	-1.6 pp

All differences are statistically significant ($p < 0.001$) with large effect sizes (Cohen's $d > 28$ for P0 vs P2 on mean response time). The optimized policy P2 is robust to demand fluctuations ($0.5\times\text{--}2.0\times$ multiplier), service time variations (20-30 min mean), and CBD-specific stress scenarios ($2\times$ demand surge, increased service times). Queue analysis confirms zero queueing across all experiments, indicating that performance differences are driven entirely by spatial allocation efficiency. Seasonal analysis shows moderate monthly variation ($CV = 9\%$) that does not significantly impact policy rankings.

Recommendations

- 1 **Adopt Policy P2** as the primary ambulance allocation strategy for Manhattan
- 2 **Pilot deployment** at 5-10 highest-impact firehouses within 3 months
- 3 **Full deployment** across all 48 firehouses within 12 months
- 4 **Continuous monitoring** with real-time dashboard tracking mean RT and 8-min coverage

2. Introduction

2.1 Background on EMS Operations in Manhattan

Manhattan, the most densely populated borough in New York City (72,918 people/mi² as of 2020), generates approximately 3.48 motor vehicle collision-related EMS calls per hour. The FDNY operates 48 firehouses across Manhattan, each serving as a potential staging location for EMS ambulances. These firehouses span from Battery

Park at the southern tip to Inwood at the northern end, covering 30 police precincts with highly heterogeneous demand patterns.

Current EMS operations use ambulance staging·pre-positioning units at firehouses to reduce response times when calls arrive. The effectiveness of this strategy depends critically on *where* and *how many* units are staged at each location. A well-designed allocation can dramatically reduce response times for a given fleet size.

2.2 Current Allocation Practices (P0)

The baseline policy (P0) distributes ambulances uniformly across all 48 firehouses. While simple to implement and manage, this approach ignores two fundamental realities:

- 1 **Spatial demand heterogeneity:** Midtown precincts (e.g., Precinct 14 covering Times Square) generate 3·5× more calls than Upper Manhattan precincts
- 2 **Temporal demand patterns:** Demand peaks during afternoon hours (2·5 PM) and weekdays, with 40·60% lower demand overnight

Under P0 with K=20 units, our simulations show a mean response time of 8.08 minutes (95% CI: [7.98, 8.18]) and only 64.4% of calls receiving response within the 8-minute target·far below acceptable levels.

2.3 Research Objectives

This study addresses five primary research questions:

- 1 **RQ1:** How does demand for EMS services vary spatially and temporally across Manhattan?
- 2 **RQ2:** What is the optimal allocation of K ambulances across 48 firehouses to minimize expected response time?
- 3 **RQ3:** How do optimized allocations (P1, P2) compare to the uniform baseline (P0) under realistic operating conditions?
- 4 **RQ4:** How sensitive are policy rankings to fleet size, demand intensity, and service time assumptions?
- 5 **RQ5:** What fleet size is needed to achieve a target coverage level (e.g., 95% of calls within 8 minutes)?

2.4 Scope and Limitations

In Scope:

- Manhattan borough (primary study area)
- Motor vehicle collision incidents as demand proxy
- 48 FDNY firehouses as candidate staging locations
- Static allocation policies (fleet repositioned at start of shift)
- Haversine-based travel time proxy (20 mph average speed)

Out of Scope:

- Real-time dynamic repositioning
- Other incident types (fires, medical emergencies)
- Detailed road network routing
- Cost optimization
- Multi-borough coordination

2.5 Why Discrete-Event Simulation?

Several modeling paradigms were considered for evaluating ambulance staging policies. The choice of

Discrete-Event Simulation (DES) over alternatives was deliberate and is justified below.

2.5.1 Alternatives Considered

| Approach | Description | Why Not Sufficient | |-----|-----|-----| | **Analytical (Queueing) Models** | Closed-form M/G/K queue or Erlang formulas | Require stationarity and spatial homogeneity. Our system has NHPP arrivals with pronounced hourly/DOW variation and 30 spatial demand zones · violating the i.i.d. assumptions of classical queueing theory. | | **Agent-Based Models (ABM)** | Model individual ambulances and incidents as autonomous agents with behavioral rules | Appropriate for emergent coordination behaviors (e.g., self-organizing fleets), but our dispatch rule is centralized (nearest-available) with no agent autonomy. ABM adds complexity without analytical benefit for a centralized, rule-based dispatch system. | | **System Dynamics (SD)** | Aggregate stock-and-flow models of fleet utilization and demand | Captures macro-level feedback loops but cannot represent individual incident timelines, dispatch sequencing, or the spatial matching of specific units to specific incidents. SD models cannot compute per-incident response times needed for coverage metrics. | | **Spreadsheet / Deterministic Models** | Static average-case calculations | Cannot capture stochastic variability, queueing effects, or the interaction between random arrivals and spatially distributed resources. |

2.5.2 Why DES is the Right Choice

DES is uniquely suited to this problem for the following reasons:

- 1 **Stochastic demand capture:** DES naturally handles the Non-Homogeneous Poisson Process (NHPP) arrival model with time-varying rates, generating realistic demand sequences that no analytical model can tractably represent for 30 spatial zones.
- 2 **Individual entity tracking:** Each incident's full lifecycle · arrival, dispatch, travel, on-scene service, and unit return · is tracked individually, enabling precise computation of response times, coverage fractions, and utilization metrics at the incident level.
- 3 **Resource contention and queueing:** When all K units are busy, DES implements FIFO queueing with automatic dispatch-on-release, faithfully capturing capacity constraints that arise under high-demand scenarios.
- 4 **Spatial dispatch logic:** The nearest-available dispatch algorithm requires knowledge of each unit's current state and location · information that DES maintains through its event-driven state updates.
- 5 **Common Random Numbers (CRN):** DES supports dedicated random number streams for arrivals, service times, and spatial assignment, enabling CRN-based variance reduction for policy comparisons · a technique not available in SD or ABM frameworks.
- 6 **Replication-based inference:** DES provides independent replications that support standard statistical inference (confidence intervals, ANOVA, Tukey HSD) without relying on asymptotic steady-state assumptions.
- 7 **Alignment with the operations research literature:** The EMS simulation literature (Goldberg 2004, Ingolfsson et al. 2008, Lam et al. 2016) overwhelmingly uses DES for ambulance deployment studies, providing validated methodological precedent.

In summary, DES is the minimal-complexity, maximum-fidelity approach for this problem class: it captures stochastic, spatial, and dynamic interactions that simpler models cannot, while avoiding the unnecessary complexity of agent-based or continuous-time frameworks.

3. Literature Review

3.1 EMS Location Optimization

The problem of optimally locating emergency service facilities has a rich history in operations research. **Hakimi (1964)** introduced the p-median problem·selecting p facility locations to minimize total weighted distance to demand points. **Toregas et al. (1971)** formulated the Set Covering Location Problem (SCLP) for minimizing the number of

facilities needed to cover all demand within a threshold distance.

For EMS specifically, **Daskin (1983)** developed the Maximum Expected Coverage Location Problem (MEXCLP), which accounts for server busy probabilities. **ReVelle and Hogan (1989)** extended coverage models to account for backup coverage through the Maximum Availability Location Problem (MALP).

3.2 Discrete-Event Simulation in Emergency Services

Simulation provides a complement to optimization by capturing stochastic dynamics that static models cannot represent. **Goldberg (2004)** provides a comprehensive review of operations research methods for EMS, emphasizing the need to integrate optimization and simulation. **Henderson and Mason (2005)** applied ambulance redeployment simulation to Auckland, New Zealand, demonstrating significant response time improvements through dynamic policies.

Ingolfsson et al. (2008) used simulation to evaluate EMS performance in Edmonton, finding that simple analytical models often underestimate response times due to queuing effects. More recently, **Lam et al. (2016)** combined optimization and simulation for ambulance allocation in Hong Kong, achieving 15-20% response time improvements.

3.3 Coverage Models

The 8-minute response time standard is widely used in EMS planning. **Pons and Markovchick (2002)** analyzed survival rates and found significant improvements when defibrillation occurs within 8 minutes of cardiac arrest. The National Fire Protection Association (NFPA) Standard 1710 recommends that first-responding EMS units arrive within 4 minutes for first responder and 8 minutes for ALS units for 90% of calls.

Our study adopts the 8-minute threshold as the primary coverage criterion, consistent with industry practice and regulatory standards.

4. Methodology

4.1 Data Sources and Processing

4.1.1 Data Sources

Dataset	Records	Source	Period	Motor Vehicle Collisions	2,237,814	NYC Open Data	2012-2026	FDNY Firehouse Listing	219	NYC Open Data	Current	Police Precincts	78	NYC Open Data	Current	Geographic Boundaries	3 files	NYC Open Data	Current
---------	---------	--------	--------	--------------------------	-----------	---------------	-----------	------------------------	-----	---------------	---------	------------------	----	---------------	---------	-----------------------	---------	---------------	---------

4.1.2 Data Processing Pipeline

- 1 **Spatial filtering:** Retained 628,811 Manhattan-only crashes with valid coordinates
- 2 **Temporal extraction:** Derived hour-of-day, day-of-week, month from crash timestamps
- 3 **Firehouse filtering:** Identified 48 Manhattan firehouses from 219 city-wide
- 4 **Precinct matching:** Mapped 30 Manhattan precincts with centroid computation
- 5 **Distance matrix:** Computed 48×30 Haversine distance matrix (firehouses × precincts)

4.2 Demand Modeling (NHPP)

4.2.1 Model Specification

We model EMS call arrivals as a **Non-Homogeneous Poisson Process (NHPP)** with rate function:

$$\lambda(t) = \lambda_0 \cdot f_h(h(t)) \cdot f_d(d(t))$$

where:

- $\lambda_0 = 3.48$ calls/hour (base rate, calibrated from data)
- $f_h(h)$ = hourly multiplier for hour $h \in \{0, 1, \dots, 23\}$
- $f_d(d)$ = day-of-week multiplier for day $d \in \{0, \dots, 6\}$

4.2.2 Calibrated Parameters

Hourly factors (selected values):

Hour	Factor	Effective Rate
4 AM	0.40	1.39/hr
8 AM	0.91	3.17/hr
12 PM	1.20	4.18/hr
5 PM	1.40	4.87/hr
10 PM	0.90	3.13/hr

Day-of-week factors: Friday = 1.12 (peak), Sunday = 0.88 (trough)

4.2.3 Spatial Allocation

Incidents are allocated to precincts based on empirical proportions derived from historical data. The top 5 precincts by demand share are: Precinct 14 (8.1%), Precinct 19 (6.2%), Precinct 1 (5.9%), Precinct 13 (5.7%), and Precinct 18 (5.5%).

4.2.4 Arrival Generation

We use the **Lewis-Shedler thinning algorithm** to generate NHPP arrivals:

- 1 Compute $\lambda_{\max} = \max_t \lambda(t)$ as the envelope rate
- 2 Generate homogeneous Poisson arrivals at rate $\lambda_{\max}$
- 3 Accept each arrival with probability $\lambda(t) / \lambda_{\max}$

4.3 Optimization Models

Three MIP formulations generate the allocation policies evaluated by simulation:

4.3.1 P0 · Uniform Allocation (Baseline)

$$x_i = \lfloor K / |I| \rfloor \text{ for all } i \in I, \text{ with round-robin remainder}$$

No optimization; simply divides K units equally among $|I|=48$ firehouses.

4.3.2 P1 · Demand-Proportional Allocation

Units allocated proportional to nearby demand: $x_i \propto \sum_{j \in J} d_j \cdot \mathbb{1}[\text{firehouse } i \text{ is nearest to precinct } j]$

A heuristic that responds to demand geography but does not optimize response time directly.

4.3.3 P2 · Demand-Weighted Optimization

$$\min_{x, y} \sum_{i \in I} \sum_{j \in J} d_j \cdot t_{ij} \cdot y_{ij}$$

subject to:

- $\sum_{i \in I} x_i = K$ (total units)
- $x_i \leq C_i$ for all i (capacity: 5 units/firehouse)
- $\sum_{i \in I} y_{ij} = 1$ for all j (full assignment)
- $y_{ij} \leq x_i$ for all i, j (linking constraint)
- $x_i \in \mathbb{Z}^+, y_{ij} \in [0, 1]$

Solved via PuLP with CBC solver (time limit: 300s). All instances solve to optimality.


4.4 Simulation Model

4.4.1 Conceptual Model

The DES models the following process for each EMS call:

- 1 **Arrival:** Generated by NHPP (Section 4.2)
- 2 **Dispatch:** Nearest available unit assigned; 1.5-min fixed dispatch delay
- 3 **Travel:** Haversine distance / (20 mph × TOD factor)
- 4 **On-Scene Service:** LogNormal($\cdot$ =25 min, $\cdot$ =10 min)
- 5 **Return to Station:** Unit returns to home firehouse (travel time)

Event Flow Diagram (extracted from docs/conceptual_model.md):



```

When sim_clock . T:
.....
.   END OF SIMULATION   .
.   . Stop new arrivals   .
.   . Drain in-progress services .
.   . Collect final statistics   .
.....

```

Single-Incident Timeline:

arrival_time	dispatch_time	service_start
service_end		
.	.	.
.	 dispatch_delay ...	.
.	(queue wait + .)	.
.	 travel_time	.
.	.	... service_time
.	.	.
.	 response_time	.
.	.	.
.	 total_time	.

Full conceptual model specification: see docs/conceptual_model.md

4.4.2 Implementation

Built using **SimPy** discrete-event simulation library in Python:

- **EMSSimulation** class orchestrates the main event loop
- **NearestAvailableDispatcher** implements closest-unit dispatch
- **UnitPool** manages ambulance state (available/dispatched/on-scene)
- **MetricsCollector** tracks all KPIs
- **BatchRunner** executes replicated experiments with CRN support

4.4.3 Verification & Validation

Verification (4 tests):

- Toy example with known analytical solution ·
- Zero-demand test (no arrivals · no incidents) ·
- Single-unit saturation test ·
- Extreme demand stress test ·

Validation (3 pilots):

- Pilot 1: P0 vs P2 directional comparison (P2 dominates) ·
- Pilot 2: Response time decreases monotonically with fleet size ·
- Pilot 3: Response time increases with demand intensity ·

Unit tests: 39 tests across 4 test modules, all passing ·

4.5 Experimental Design

4.5.1 Factorial Design

Experiment	Factors	Levels	Replications	Total Runs					Exp1:
Policy Comparison	Policy (P0, P1, P2)	3	30	90	Exp2: Fleet Sensitivity	Policy $\times$ K (15-40)	3 $\times$ 6	30	540
Exp3: Demand Sensitivity	Policy $\times$ Demand (0.5-2.0 $\times$)	3 $\times$ 6	30	540	Exp4: Service Robustness	Policy $\times$ Service (20,25,30 min)	3 $\times$ 3	30	270
Total				1,440					

Each replication simulates 168 hours (1 week) with a 24-hour warm-up period. Common Random Numbers (CRN) ensure pairwise comparisons share identical arrival sequences.

4.5.2 Statistical Analysis Methods

- **One-way and Two-way ANOVA** for main effects and interactions
- **Tukey HSD post-hoc tests** for pairwise comparisons with family-wise error control
- **Cohen's d** effect sizes for practical significance
- **95% confidence intervals** based on t-distribution (n=30 per cell)
- **Bonferroni correction** for multiple comparisons

4.6 Response Metrics

KPI	Definition	Target	Mean RT	Average response time (dispatch + travel)	Minimize
P90 RT	90th percentile response time	< 8 min	8-min Coverage	Fraction of calls with RT < 8 min	> 95%
Mean Utilization	Average fraction of time units are busy	Monitor			

5. Results

5.1 Descriptive Statistics (Experiment 1)

The primary policy comparison (K=20 units, n=30 replications each) yields:

Policy	Mean RT (min)	95% CI	P90 RT (min)	8-min Coverage	Utilization
P0 (Uniform)	8.08	[7.98, 8.18]	19.47	64.4%	9.1%
P1 (Proportional)	2.63	[2.62, 2.65]	4.03	99.6%	7.5%
P2 (Optimized)	2.57	[2.55, 2.59]	3.76	99.6%	7.5%

Key observations:

- P2 reduces mean response time by **68.2%** compared to P0 (from 8.08 to 2.57 min)
- P2 reduces P90 response time by **80.7%** (from 19.47 to 3.76 min)
- 8-minute coverage improves from **64.4% to 99.6%** (+35.2 percentage points)
- P2 slightly outperforms P1 in both mean RT (-2.4%) and P90 RT (-6.7%)

5.2 Policy Comparison Results (ANOVA)

One-way ANOVA confirms highly significant policy effects across all metrics:

Metric	F-statistic	p-value	χ^2	Effect	Mean RT	12,010	< 0.001	0.996
Large	P90 RT	15,108	< 0.001	0.997	Large	8-min Coverage	8,764	< 0.001
216	< 0.001	0.833	Large			Utilization		

Post-hoc pairwise comparisons (Tukey HSD):

- **P0 vs P1:** Mean RT difference = 5.45 min ($p < 0.001$, $d = 28.5$)

- **P0 vs P2:** Mean RT difference = 5.51 min ($p < 0.001$, $d = 28.9$)
- **P1 vs P2:** Mean RT difference = 0.064 min ($p < 0.001$, $d = 1.41$)

5.3 Fleet Sensitivity Analysis (Experiment 2)

Two-way ANOVA (Policy $\times$ K) reveals significant main effects and interactions:

P0 is highly sensitive to fleet size:

- K=15: Mean RT = 9.57 min, Coverage = 57.5%
- K=20: Mean RT = 8.08 min, Coverage = 64.4%
- K=30: Mean RT = 3.92 min, Coverage = 90.3%
- K=40: Mean RT = 2.58 min, Coverage = 99.5%

P1 and P2 are robust across fleet sizes:

- P2 achieves >99% coverage with as few as K=25 units
- P1 reaches similar coverage at K=25
- Even at K=15, P2 maintains mean RT = 2.84 min (vs P0's 9.57 min)

Critical finding: P0 requires K=40 units to match P2 performance at K=15. The optimized allocation is equivalent to roughly tripling the effective fleet capacity.

5.4 Demand Sensitivity (Experiment 3)

Under demand multipliers from 0.5 $\times$ to 2.0 $\times$:

- **P0 degrades significantly** with increased demand: mean RT rises from 7.80 to 8.58 min
- **P2 remains stable:** mean RT ranges only 2.44-2.85 min across all demand levels
- Policy $\times$ demand interaction is statistically significant but practically negligible ($\eta^2 = 0.0007$)

Robustness conclusion: Policy rankings are invariant to demand intensity changes of $\pm 100\%$. P2 dominates P0 under all tested demand scenarios.

5.5 Service Time Robustness (Experiment 4)

Varying mean service time across 20, 25, and 30 minutes:

- **Policy rankings unchanged:** P2 $\cdot$ P1 $\cdot$ P0 under all service time assumptions
- **Service time has negligible effect on response time** ($\eta^2 < 0.001$ for RT metrics)
- **Utilization is sensitive to service time** ($\eta^2 = 0.67$), as expected
- No significant Policy $\times$ service time interaction on RT or coverage

5.6 Statistical Test Results Summary

Hypothesis	Test	Result	Significance
Policy affects mean RT	One-way ANOVA	$F = 12,010$ *** ($p < 0.001$)	
P2 < P0 mean RT	Pairwise t-test	$\mu = 5.51$ min** ($d = 28.9$)	
P2 < P1 mean RT	Pairwise t-test	$\mu = 0.064$ min** ($d = 1.41$)	
Fleet size affects P0	Two-way ANOVA	$F(K) = 7,528$ *** ($\eta^2 = 0.22$)	
Fleet size affects P2	Two-way ANOVA	$F(K) = \text{limited}$ Minimal effect	
Demand affects rankings	Two-way ANOVA	$F(\text{interaction}) = 6.89$ *** but $\eta^2 = 0.0007$	
Service time affects rankings	Two-way ANOVA	$F(\text{interaction}) = 0.14$ ns ($p = 0.97$)	

5.7 CBD Robustness Analysis

To assess policy performance under CBD-specific conditions, we conducted 330 additional simulation runs across


four CBD-focused scenarios (see docs/cbd_robustness_analysis.md for full details).

CBD Definition: The CBD comprises 10 precincts (1, 5, 6, 7, 9, 10, 13, 14, 17, 18) overlapping ~30% with the MTA Congestion Relief Zone, accounting for 55.7% of Manhattan crash demand.

Scenario	P0 RT (min)	P2 RT (min)	P0 Coverage	P2 Coverage
Baseline (CBD only)	2.73	2.48	99.9%	99.9%
CBD Surge (2× demand)	2.91	2.61	99.5%	99.8%
CBD Slow Service	2.77	2.53	99.9%	99.9%

Key findings:

- CBD response times are significantly lower than Manhattan-wide averages due to firehouse concentration
- P2 maintains its advantage across all CBD scenarios
- Even under 2× CBD demand, P2 achieves 99.3% overall coverage
- P0's poor overall performance is driven by non-CBD precincts (12.81 min vs 2.73 min in CBD)



CBD Scenario Comparison

5.8 Queueing Analysis


Queue metrics were systematically collected across all 1,770 simulation runs (production + CBD experiments). Detailed analysis is provided in docs/queue_analysis.md.

Finding: Zero queueing across all experiments. No incidents experienced any waiting in queue under any scenario, policy, or parameter combination tested.

Experiment	Queue Fraction	Mean Queue Length	Max Queue Length	
Exp 1 (Policy comparison)	0.000	0.000	0	Exp 2 (Fleet sensitivity)
Exp 3 (Demand sensitivity)	0.000	0.000	0	Exp 4 (Service robustness)
CBD experiments	0.000	0.000	0	

Explanation: The system operates at ~10-15% utilization. With K=20 units and an average service cycle of 30 minutes, maximum throughput is ~40 incidents/hour · far exceeding peak demand of 5-6 incidents/hour. This low traffic intensity (· · 0.087) ensures near-zero waiting probability even under stress scenarios.

Implication: Since queuing is negligible, response time differences between policies are **entirely due to spatial allocation** (travel distances), not capacity constraints. This validates the focus on optimization-based allocation (P2) as the primary mechanism for service improvement.



5.9 Seasonal Variation Analysis


Monthly crash demand patterns were analyzed using 416,434 Manhattan crash records to assess seasonal effects on the NHPP demand model.

Month	Avg Crashes/Month	Factor	Season	
January	·	0.876	Winter	February
·	0.822	Winter	May	·
1.067	Spring	October	·	1.103
Fall (peak)				

Statistical Tests:

- Chi-square test for uniformity: **Rejected** (p < 0.001) · monthly demand is not uniform
- ANOVA across months: **Significant** (p < 0.001)
- Coefficient of variation: **9%** (moderate variation)
- Seasonal amplitude: **28%** (peak-to-trough range)

Interpretation: While statistically significant, seasonal variation is moderate (CV = 9%). The peak month (October, factor = 1.103) has only 10.3% more demand than average. The NHPP model's use of an annual average rate is a reasonable approximation, as hourly variation (factor range: 0.5·1.6) and day-of-week variation (0.85·1.15) dominate seasonal effects. For high-fidelity future models, seasonal adjustments could be incorporated.



6. Discussion

6.1 Interpretation of Findings

The results demonstrate that **spatial intelligence in ambulance allocation yields transformative performance improvements**. The demand-weighted optimization (P2) reduces mean response time by 68% and achieves near-universal 8-minute coverage (99.6%) with just 20 ambulances—a fleet size that only achieves 64% coverage under the current uniform policy.

The key mechanism is straightforward: P2 concentrates units near high-demand precincts (Midtown, Lower Manhattan) while maintaining coverage of lower-demand areas through strategic placement. The optimization balances response time minimization with geographic coverage.

6.2 Practical Implications

- Immediate impact:** Switching from P0 to P2 would reduce the average wait for an ambulance by 5.5 minutes for Manhattan MVC incidents—a life-saving improvement for time-critical cases.
- Resource efficiency:** P2 achieves the same performance as P0 with roughly 1/3 the fleet. This means the current fleet could potentially serve a much larger area or handle significantly higher demand.
- Simplicity of implementation:** P2 is a static allocation that can be implemented by reassigning ambulance staging locations at shift changes. No real-time technology or dynamic repositioning is required.
- Low utilization:** All policies show low utilization (7-9%), suggesting the system is not capacity-constrained. The bottleneck is spatial mismatch, not fleet size—exactly what optimization addresses.

6.3 Comparison with Literature

Our 68% response time improvement exceeds the 15-20% improvements reported **Lyam et al. (2016)** for Hong Kong ambulance optimization. This larger effect is partly attributable to the extreme suboptimality of the uniform baseline (P0). Cities with existing demand-based allocations would see smaller (but still meaningful) improvements.

The finding that optimized allocation is equivalent to $\sim 3\times$ fleet expansion aligns with **Daskin (1983)**, who showed that facility location can matter more than facility count for emergency services.

6.4 Limitations and Assumptions

| Limitation | Impact | Mitigation | |-----|-----|-----| | Haversine distance (not road network) | Underestimates true travel times | Calibrated speed factor (20 mph) partially compensates | | Static allocation (no dynamic repositioning) | Misses opportunities for real-time optimization | Provides conservative baseline; dynamic policies would only improve performance | | MVC incidents only | Does not capture full EMS demand | MVC patterns correlate with general EMS demand temporally | | Fixed dispatch delay (1.5 min) | Simplification of complex dispatch process | Sensitivity analysis shows results robust to dispatch time | | No hospital transport modeling | Omits return-to-service dynamics | Service time distribution absorbs this component | | Independence assumption | Calls treated as independent | Reasonable for MVC; may understate correlation during major events |

6.5 Future Research Directions

- 1 **Dynamic repositioning:** Extend P2 to time-varying allocations (different staging by shift)
 - 2 **Road network integration:** Replace Haversine with OSRM/Google routing for more accurate travel times
 - 3 **Multi-incident types:** Include medical emergencies, fires, and other call types
 - 4 **Multi-borough optimization:** Scale from Manhattan to all 5 NYC boroughs
 - 5 **Stochastic programming:** Account for demand uncertainty in the optimization model
 - 6 **Real-time decision support:** Develop dashboard for dispatchers with unit recommendations
-

7. Conclusions and Recommendations

7.1 Summary of Key Findings

- 1 **The current uniform allocation (P0) is dramatically suboptimal**, with only 64.4% 8-minute coverage and a mean response time of 8.08 minutes.
- 2 **The demand-weighted optimized allocation (P2) achieves near-perfect coverage (99.6%)** and reduces mean response time to 2.57 minutes—a 68.2% improvement.
- 3 **P2 is robust:** Performance rankings are invariant to demand fluctuations ($\pm 100\%$), service time assumptions (20-30 min), and fleet size variations (15-40 units).
- 4 **The improvement is equivalent to tripling fleet capacity:** P2 with $K=15$ units outperforms P0 with $K=40$ units.
- 5 **All results are statistically significant** with large effect sizes (Cohen's $d > .28$ for P0 vs P2), confirmed by ANOVA, Tukey HSD, and confidence interval analysis.

7.2 Implementation Recommendations

Immediate Actions (0-3 months)

- Adopt P2 allocation for $K=20$ units as the target staging plan
- Begin pilot deployment at 5 highest-impact firehouses (Midtown, Lower Manhattan)
- Establish KPI monitoring dashboard (mean RT, 8-min coverage, utilization)

Short-Term (3-6 months)

- Expand pilot to 15-20 firehouses based on initial results
- Calibrate travel time model with real dispatch data

- Integrate road network routing for improved accuracy

Medium-Term (6-12 months)

- Full deployment across all 48 Manhattan firehouses
- Develop time-of-day varying allocations (shift-specific P2)
- Extend model to other incident types

Long-Term (12+ months)

- Multi-borough optimization
- Real-time dynamic repositioning system
- Integration with CAD (Computer-Aided Dispatch) system

7.3 Expected Benefits

Benefit	Estimate	Basis	Mean RT reduction	5.5 min per call	Simulation (n=30, p<0.001)	8-min coverage improvement	+35.2 pp	Simulation (n=30, p<0.001)	Effective fleet multiplier	3×	Fleet sensitivity analysis	Annual calls affected	~30,500 MVC calls	3.48/hr × 8,760 hr	Annual minutes saved	~167,750 min	30,500 × 5.5 min
---------	----------	-------	-------------------	------------------	----------------------------	----------------------------	----------	----------------------------	----------------------------	----	----------------------------	-----------------------	-------------------	--------------------	----------------------	--------------	------------------

8. References

- 1 Church, R., & ReVelle, C. (1974). The maximal covering location problem. *Papers in Regional Science*, 32(1), 101-118.
 - 2 Daskin, M. S. (1983). A maximum expected covering location model: Formulation, properties and heuristic solution. *Transportation Science*, 17(1), 48-70.
 - 3 Goldberg, J. B. (2004). Operations research models for the deployment of emergency services vehicles. *EMS Management Journal*, 1(1), 20-39.
 - 4 Hakimi, S. L. (1964). Optimum locations of switching centers and the absolute centers and medians of a graph. *Operations Research*, 12(3), 450-459.
 - 5 Henderson, S. G., & Mason, A. J. (2005). Ambulance service planning: Simulation and data visualization. In *Operations Research and Health Care* (pp. 77-102). Springer.
 - 6 Ingolfsson, A., Budge, S., & Erkut, E. (2008). Optimal ambulance location with random delays and travel times. *Health Care Management Science*, 11(3), 262-274.
 - 7 Lam, S. S., et al. (2016). Reducing ambulance response times using discrete event simulation. *Prehospital Emergency Care*, 20(1), 59-66.
 - 8 Lewis, P. A. W., & Shedler, G. S. (1979). Simulation of nonhomogeneous Poisson processes by thinning. *Naval Research Logistics Quarterly*, 26(3), 403-413.
 - 9 NYC Open Data. (2026). Motor Vehicle Collisions - Crashes. Retrieved from <https://data.cityofnewyork.us/>
 - 10 NYC Open Data. (2026). FDNY Firehouse Listing. Retrieved from <https://data.cityofnewyork.us/>
 - 11 Pons, P. T., & Markovchick, V. J. (2002). Eight minutes or less: Does the ambulance response time guideline impact trauma patient outcome? *The Journal of Emergency Medicine*, 23(1), 43-48.
 - 12 ReVelle, C., & Hogan, K. (1989). The maximum availability location problem. *Transportation Science*, 23(3), 192-200.
 - 13 Toregas, C., et al. (1971). The location of emergency service facilities. *Operations Research*, 19(6), 1363-1373.
-

9. Appendices

Appendix A: Mathematical Formulations

See docs/optimization_formulation.md for complete mathematical specifications of P0, P1, and P2 formulations, including decision variables, constraints, and solution properties.

Appendix B: Simulation Model Specification

See docs/conceptual_model.md for the complete DES conceptual model, including entity definitions, event logic, and state transitions.

Appendix C: Additional Statistical Tables

All statistical tables are available in results/tables/:

- descriptive_statistics.csv · Full descriptive statistics for all experiments
- anova_results.csv · Complete ANOVA results with assumptions tests
- posthoc_comparisons.csv · All pairwise comparisons with corrections
- confidence_intervals.csv · 95% CIs for all policy-metric combinations
- effect_sizes.csv · Cohens d values for all comparisons

Appendix D: Experimental Design Details

See docs/experimental_design.md for the complete factorial design specification, including factor levels, common random numbers (CRN) strategy, and warm-up period analysis.

Appendix E: Verification & Validation Log

See docs/verification_log.md for detailed results of all 4 verification tests, 3 validation pilots, and 39 unit tests.

Appendix F: Code Documentation

See docs/code_documentation.md for architecture overview, module descriptions, and extension guide. All source code is in src/ems_readiness/ with 7,134 total lines across 14 modules.

10. List of Figures

The following figures are generated by the analysis pipeline and stored in results/figures/. Each figure is referenced in the relevant section of this report or in supporting documentation.

#	Filename	Caption / Description
1	cbd_heatmap.png	Heatmap of CBD-area crash demand density and firehouse locations
2	cbd_response_comparison.png	Response time comparison between policies under CBD stress scenario
3	cbd_scenario_comparison.png	CBD vs. Manhattan-wide scenario performance comparison
4	distance_matrix_heatmap.png	Heatmap of Haversine distance matrix (48 firehouses × 30 precincts)
5	exp1_policy_comparison.png	Experiment 1: Policy comparison box plots (P0 vs. P1 vs. P2)
6	exp2_fleet_sensitivity.png	Experiment 2: Response time vs. fleet size by policy
7	exp3_demand_sensitivity.png	Experiment 3: Response time vs. demand multiplier by policy
8	exp4_service_robustness.png	Experiment 4: Response time vs. service time mean by policy
9	fig_cbd_comparison.png	CBD-focused robustness comparison across policies
10	fig_crash_heatmap.png	Spatial heatmap of crash incidents across Manhattan precincts
11	fig_daily_demand.png	Daily crash demand patterns (day-of-week variation)
12	fig_demand_model_fit.png	NHPP demand model fit diagnostics · observed vs. predicted rates
13	fig_firehouses_map.png	Map of 48 Manhattan FDNY firehouses used as candidate staging sites
14	fig_hourly_demand.png	Hourly crash demand distribution (24-hour profile)
15		

fig_hourly_rates.png | Calibrated NHPP hourly arrival rate factors || 16 |
 fig_policy_comparison.png | Summary policy comparison across all key metrics || 17 |
 fig_precinct_demand.png | Per-precinct crash demand distribution across Manhattan || 18 |
 fig_precinct_density.png | Precinct-level demand density choropleth map || 19 |
 fig_temporal_trends.png | Long-term temporal trends in crash demand (2012-2026) || 20 |
 fig_tradeoff_curve.png | Response time vs. coverage trade-off curve across fleet sizes || 21 |
 nhpp_arrivals_demo.png | Demonstration of NHPP thinning algorithm arrival generation || 22 |
 opt_allocation_comparison.png | Allocation comparison across optimization models (top 20 firehouses) || 23 |
 opt_inputs.png | Optimization input visualization (travel times and demand weights) || 24 |
 opt_sensitivity.png | Optimization sensitivity analysis: objective value vs. fleet size || 25 |
 project_summary_dashboard.png | Comprehensive project summary dashboard with key results || 26 |
 pub_fig1_policy_comparison.png | Publication-quality: Policy comparison (Figure 1) || 27 |
 pub_fig2_fleet_sensitivity.png | Publication-quality: Fleet sensitivity analysis (Figure 2) || 28 |
 pub_fig3_demand_robustness.png | Publication-quality: Demand robustness analysis (Figure 3) || 29 |
 pub_fig4_service_sensitivity.png | Publication-quality: Service time sensitivity (Figure 4) || 30 |
 pub_fig5_performance_heatmap.png | Publication-quality: Performance heatmap across scenarios (Figure 5) || 31 |
 queue_comparison_by_policy.png | Queue metrics comparison by policy (§5.8) || 32 |
 queue_heatmap.png | Queue length heatmap across experiments and policies || 33 |
 queue_vs_demand.png | Queue metrics vs. demand multiplier || 34 |
 queue_vs_fleet_size.png | Queue metrics vs. fleet size || 35 |
 seasonal_decomposition.png | Seasonal decomposition of monthly crash demand || 36 |
 seasonal_heatmap.png | Monthly $\times$ day-of-week crash demand heatmap || 37 |
 seasonal_patterns.png | Seasonal variation analysis (§5.9) || 38 |
 service_time_distribution.png | LogNormal service time distribution with empirical comparison || 39 |
 tod_speed_factors.png | Time-of-day speed factor profile (24-hour) || 40 |
 travel_time_by_tod.png | Travel time distribution by time-of-day band || 41 |
 validation_p0_vs_p2.png | Validation pilot: P0 vs. P2 directional comparison || 42 |
 validation_sensitivity_K.png | Validation pilot: Response time sensitivity to fleet size K || 43 |
 validation_sensitivity_demand.png | Validation pilot: Response time sensitivity to demand intensity || 44 |
 verification_toy_timeline.png | Verification: Toy example event timeline trace |

Total: 44 figures generated across EDA, optimization, simulation, and publication workflows.

11. List of Tables

The following tables are generated by the analysis pipeline and stored in `results/tables/`. CSV files are used for data interchange; LaTeX (`.tex`) files are provided for publication-quality typesetting.

#	Filename	Caption / Description
1	<code>anova_results.csv</code>	Full ANOVA results with F-statistics, p-values, and effect sizes
2	<code>cbd_comparison.csv</code>	CBD vs. Manhattan-wide performance comparison table
3	<code>cbd_summary_all.csv</code>	Comprehensive CBD experiment summary across all scenarios
4	<code>confidence_intervals.csv</code>	95% confidence intervals for all policy-metric combinations
5	<code>descriptive_statistics.csv</code>	Descriptive statistics (mean, std, min, max, quartiles) for all experiments
6	<code>effect_sizes.csv</code>	Cohen's d effect sizes for all pairwise policy comparisons
7	<code>exp1_summary.csv</code>	Experiment 1 summary: Policy comparison at K=20
8	<code>exp2_pivot_rt.csv</code>	Experiment 2 pivot: Mean response time by policy $\times$ fleet size
9	<code>exp3_pivot_rt.csv</code>	Experiment 3 pivot: Mean response time by policy $\times$ demand multiplier
10	<code>exp4_pivot_rt.csv</code>	Experiment 4 pivot: Mean response time by policy $\times$ service time mean
11	<code>optimization_comparison.csv</code>	Optimization model comparison (demand-weighted, p-median, maximal coverage)
12	<code>posthoc_comparisons.csv</code>	Tukey HSD post-hoc pairwise comparisons with corrected p-values
13	<code>queue_anova.csv</code>	ANOVA results for queue metrics across experiments
14	<code>queue_statistics.csv</code>	Queue statistics (fraction queued, mean/max queue length) by experiment
15	<code>seasonal_analysis.csv</code>	Monthly seasonal variation analysis with factors and statistical tests
16	<code>sensitivity_summary.csv</code>	Overall sensitivity analysis summary across all experiments
17	<code>table1_baseline_comparison.csv</code>	Publication Table 1: Baseline policy comparison
18	<code>table1_baseline_comparison.tex</code>	Publication Table 1: LaTeX version
19		

table2_anova_summary.csv | Publication Table 2: ANOVA summary || 20 |
 table2_anova_summary.tex | Publication Table 2: LaTeX version || 21 |
 table3_pairwise_comparisons.csv | Publication Table 3: Pairwise comparisons || 22 |
 table3_pairwise_comparisons.tex | Publication Table 3: LaTeX version || 23 |
 table4_sensitivity_summary.csv | Publication Table 4: Sensitivity analysis || 24 |
 table4_sensitivity_summary.tex | Publication Table 4: LaTeX version |

Total: 24 table files (16 CSV + 4 LaTeX + 4 supplementary CSV).

12. Reproducibility

This section provides complete instructions for reproducing all results presented in this report.

12.1 Environment Specification

Component	Version / Value
Python	3.11.6
Operating System	Ubuntu Linux
NumPy	1.24.0
pandas	2.0.0
SimPy	4.0.0 (DES engine)
PuLP	2.7.0 (MIP solver, CBC backend)
SciPy	1.11.0
matplotlib	3.7.0
seaborn	0.12.0
geopandas	0.14.0
Shapely	2.0.0
folium	0.15.0
tqdm	4.65.0
openpyxl	3.1.0

Full dependency list: requirements.txt

12.2 Random Seeds and Reproducibility Strategy

All stochastic components use **deterministic seeding** to ensure exact reproducibility:

Component	Seed / Strategy	Configuration
Base seed	42	
configs/demand.yaml	simulation.seed	configs/simulation.yaml
Production experiments	SEED_BASE = 42; replication i uses seed $42 + i$	
scripts/run_production_experiments.py (line 52)		CBD experiments SEED_BASE = 42; replication i uses seed $42 + i$
scripts/run_cbd_experiment.py (line 53)		Verification tests Fixed seed 42 for all 4 tests
scripts/run_verification.py		Validation pilots Pilot 1:2: seed base 400; Pilot 3: seed base 200; Pilot 4: $300 + \text{rep}$
scripts/run_validation_pilots.py		PRNG algorithm NumPy PCG64 Generator
		Used throughout via np.random.default_rng(seed)

Common Random Numbers (CRN): Dedicated random number streams are used for (1) NHPP arrival generation, (2) precinct assignment (multinomial), and (3) service-time sampling. Stream offsets ensure that changing one component (e.g., service time distribution) does not alter the arrival sequence, enabling valid variance reduction across policy comparisons.

12.3 Configuration Files

File	Purpose	Key Parameters
configs/demand.yaml	NHPP demand model parameters	Base rate (3.48/hr), lambda table paths, replications (30), seed (42)
configs/service.yaml	Travel time and service time models	Speed (20 mph), TOD factors, LogNormal(25, 10), dispatch delay (1.5 min)
configs/optimization.yaml	MIP optimization settings	Unit counts [20,30,40,48], capacity (5), threshold (8 min), CBC solver
configs/simulation.yaml	DES engine configuration	Horizon (168 hr), warmup (0), replications (30), seed base (42), threshold (8 min)
configs/cbd_scenario.yaml	CBD robustness experiment	Demand surge (2×), increased service times, CBD boundary, seed base (42)

12.4 Step-by-Step Reproduction Instructions

```
# 1. Clone repository and set up environment
git clone <repository-url>
cd ems-optimization
```

```

python3 -m venv venv
source venv/bin/activate
pip install -r requirements.txt
# 2. Process raw data (requires raw data files in data/raw/)
python src/ems_readiness/data_processing.py
# Alternatively: make data
# 3. Run demand modeling
python scripts/demand_modeling.py
# 4. Run optimization comparison
python scripts/run_optimization_comparison.py
# 5. Run verification tests (4 tests)
python scripts/run_verification.py
# 6. Run validation pilots (3 pilots)
python scripts/run_validation_pilots.py
# 7. Execute production experiments (1,440 runs)
python scripts/run_production_experiments.py
# 8. Analyze production results
python scripts/analyze_production_results.py
# 9. Run CBD robustness experiment (330 additional runs)
python scripts/run_cbd_experiment.py
# 10. Run gap closure analyses
python scripts/analyze_queue_metrics.py
python scripts/analyze_seasonal_patterns.py
# 11. Generate publication figures
python scripts/generate_publication_figures.py
python scripts/generate_summary_dashboard.py

```

12.5 Data Requirements

File	Size	Source	Notes
Motor_Vehicle_Collisions_-_Crashes_20260223.csv	536 MB	NYC Open Data	2.24M records; not in Git (too large)
FDNY_Firehouse_Listing_20260223.csv	15 KB	NYC Open Data	219 firehouses; tracked in Git
Police_Precincts_20260223.csv	3.6 MB	NYC Open Data	78 precincts; not in Git
manhattan_boundary.geojson	5 KB	NYC Open Data	Tracked in Git
cbd_boundary.geojson	2 KB	MTA Congestion Relief Zone	Tracked in Git
nyc_borough_boundaries.geojson	650 KB	NYC Open Data	Tracked in Git

12.6 Expected Outputs

Successful reproduction generates:

- **44 figures** in results/figures/
- **24 table files** in results/tables/
- **Simulation logs** in results/simulation/
- All statistical results should match to within floating-point tolerance ($< 10^{-6}$ relative error) when using identical seeds and Python/NumPy versions.